

# Nuthan Reddy Gantla

nuthan.reddy999@gmail.com | +91-9474985552

NuthanReddy   
nuthanreddy 

Highly motivated Software Engineer and avid learner with 10+ yrs. experience in Data and Software Engineering.

## Professional Experience

Company	Designation / Position	Duration
Microsoft	Senior Software Engineer	Mar 2022 - Present
Walmart	Software Engineer III	Oct 2020 - Mar 2022
Sigmoid	Software Development Engineer I	Jan 2019 - Oct 2020
Capgemini	Consultant P2	Jun 2015 - Jan 2019

## Educational History

Degree / Specialization	Institute / College / School	Duration
B. Tech: Industrial & Systems Engineering	Indian Institute of Technology Kharagpur	2011 - 2015
XII: Board of Intermediate Education	Sri Chaitanya Jr College, Gudavalli	2009 - 2011
X: Board of Secondary Education	Gowtham Concept School, Gudivada	2008 - 2009

## Key Skills:

**Open Source Software:** Apache Spark, Hive, Kafka, Flink, Zookeeper, Kusto, Airflow, Delta Lake

**Other Tech Stack:** Kubernetes, Azure Synapse, ADF, PubSub, BigQuery, DataFlow, DataProc

**Databases / FS:** HDFS, S3, GCS, MongoDB, MySQL, HBase, Redis, ES, Splunk, Lake House

**Languages:** Python, Java, Scala, C#, SQL, Shell Script, PowerShell, Golang

## Projects / POCs

Project Info		Roles and Responsibilities
<b>Name</b>	<b>Azure Data Governance</b>	
Company	Microsoft	
Description	This project deals with Asset Discovery, Signal processing, Auditing and Reporting	<ul style="list-style-type: none"><li>◆ Leading <b>Next Gen Privacy - Signals Team</b> on Delete, Export Signal Delivery, Processing and Reporting.</li><li>◆ Made <b>Data Scan</b> 2-20x more performant and consistent by removing bottlenecks and non-deterministic code.</li><li>◆ Developed APIs to tag / update assets for <b>Data Map</b>.</li><li>◆ Made <b>Data Map</b> scalable with a variety of techniques.</li><li>◆ Developed deployment, monitoring and alerting frameworks for <b>Signals Team</b>.</li></ul>
Tech Stack	Python, Synapse, C#, Terraform, Kubernetes	
<b>Name</b>	<b>EBS Data Pipeline Framework</b>	
Company	Walmart	
Description	This project is about developing rule based processing engine	<ul style="list-style-type: none"><li>◆ Developed rule based data processing framework for the Enterprise Business Services Teams.</li><li>◆ Modularized and implemented CILL rules. Also made them config driven.</li><li>◆ Developed <b>Code Freeze</b> App to freeze GitHub repo from Slack using go.</li></ul>
Tech Stack	Spark, Scala, GCP, Go, Slack	
<b>Name</b>	<b>Ad Pre-bid Pipeline</b>	
Company	Sigmoid - OpenX	
Description	Scalable, Fault tolerant pipelines for faster Data Visualization.	<ul style="list-style-type: none"><li>◆ Implemented planet scale cost effective streaming pipeline to support 100k reqs/sec, 1GB/sec (14k reqs/sec/core) and 10k concurrent users which is presented in <a href="#">Apache Con NA 2020</a>.</li><li>◆ All of the services involved are deployed to GKE and tuned to scale based on the load.</li><li>◆ Implemented circuit breaker on our application to redirect data to PubSub + Dataflow to improve reliability</li></ul>
Tech Stack	Kafka, PubSub, DataFlow, Java, Helm, Kubernetes, Docker,	

		when Kafka + Kafka Connect is down keeping costs low but have high reliability.
<b>Name</b>	<b>NitroDB on BigQuery</b>	
Company	Sigmoid	
Description	BigQuery as Data Store for NitroDB	
Tech Stack	Apache Spark, BigQuery, NitroDB	<ul style="list-style-type: none"> <li>◆ Evaluated Percentile function in BigQuery and implemented <b>percentile</b> at scale using <b>t-digest</b> algorithm on <b>BigQuery</b> which can perform <b>10-100 times better</b> on pre aggregation.</li> <li>◆ Evaluated, tested and Implemented NitroDB on BigQuery which reduced costs by <b>30%</b>.</li> </ul>
<b>Name</b>	<b>NitroDB Search Optimization</b>	
Company	Sigmoid	
Description	Design search feature for SigView	
Tech Stack	Elastic Search, NitroDB	<ul style="list-style-type: none"> <li>◆ Integrated Elastic Search with NitroDB.</li> <li>◆ Developed an update/reindexing strategy.</li> </ul>
<b>Name</b>	<b>Snowflake ETL</b>	
Company	OpenX	
Description	Develop Snowflake ETL and maintain	
Tech Stack	Apache Spark, Snowflake, Python	<ul style="list-style-type: none"> <li>◆ Was one of the committers of <a href="#">spark-snowflake</a>. I fixed one of the critical bugs involving data correctness.</li> <li>◆ Developed ETL pipeline for Snowflake using Python.</li> <li>◆ Administered Snowflake for costs, permissions and External Data Source (GCS).</li> </ul>

### Conferences / Tech Blogs:

- ◆ [Apache Con NA 2020 - Google Pubsub vs Apache Kafka](#)
- ◆ [Microservices — Is it worth the hype? And How can we benefit from it?](#) - Also published by [Sigmoid](#).

### Scholastic and Personal Achievements:

- ◆ Secured an AIR 192 in AIEEE B. Arch. 2011 and AIR 376 in IIST 2011.
- ◆ Selected for *Indian National Physics Olympiad – 2011 & Indian National Chemistry Olympiad – 2011* by clearing NSEP 2010, NSEC 2010 with 99+ percentile.

### Additional Courses and Certifications:

1. Cloud Computing and Distributed Systems - IIT Patna
2. Functional Programming Principles in Scala - Coursera
3. Java Memory Management - LinkedIn Learning
4. Google Cloud Platform Big Data and Machine Learning Fundamentals - Coursera
5. Google Cloud Platform Fundamentals: Core Infrastructure - Coursera
6. Apache Kafka Series - Kafka Cluster Setup & Administration - Udemy Academy
7. Feature Selection for Machine Learning - Udemy Academy
8. AWS Certified Developer - Associate 2018 - Udemy
9. AWS Certified Solutions Architect - Associate 2018 - Udemy
10. Python for Data Science and Machine Learning Bootcamp - Udemy Academy